

## ▼ K-Nearest Neighbors (KNN) project

### Objective:

1. investigate how different values of k and distance metrics affect the accuracy of a KNN classifier on a specific dataset.
2. Visualize the decision boundaries for k iterations

### Steps:

1. Select a Dataset from UCI depository, preprocess data to handle missing values etc.
2. Implement a KNN Algorithm: for example Minkowski distance metric. (Use libraries like scikit-learn)
3. Experiment with Different k Values: Conduct experiments with various values of k and observe how the choice of k affects the performance of the model. Plot accuracy for each k value to visualize the results.
4. Experiment with Different Distance Metrics: Metrics such as Euclidean distance, Manhattan distance, and Minkowski distance with various p values (power parameter) to observe how the choice of distance metric influences the model's performance.
5. Create a synthetic two variable dataset to demonstrate decision boundaries of KNN at various iterations of k.

```
1 from google.colab import files
2 uploaded = files.upload()
```

Choose Files No file chosen  
Saving test.csv to test.csv

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

```
1 import pandas as pd
2 import numpy as np
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 from sklearn.neighbors import KNeighborsClassifier
6 from sklearn.metrics import accuracy_score
```

```
1 pip install ucimlrepo
```

```
1 from ucimlrepo import fetch_ucirepo
2 # fetch dataset
3 heart_disease = fetch_ucirepo(id=45)
4 # data (as pandas dataframes)
5 X = heart_disease.data.features
6 y = heart_disease.data.targets
```

```
1 combined_df = pd.concat([X, y], axis=1)
2 combined_df = combined_df.dropna()
3 combined_df.reset_index(drop=True, inplace=True)
```

```
1 combined_df.head(5)
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	num
0	63	1	1	145	233	1	2	150	0	2.3	3	0.0	6.0	0
1	67	1	4	160	286	0	2	108	1	1.5	2	3.0	3.0	2
2	67	1	4	120	229	0	2	129	1	2.6	2	2.0	7.0	1
3	37	1	3	130	250	0	0	187	0	3.5	3	0.0	3.0	0
4	41	0	2	130	204	0	2	172	0	1.4	1	0.0	3.0	0

```
1 from sklearn.model_selection import train_test_split
2 x= combined_df.iloc[:,0:13].values
3 y= combined_df['num'].values
4
5 x_train, x_test, y_train, y_test= train_test_split(x, y, test_size= 0.25, random_state=0)
6 from sklearn.preprocessing import StandardScaler
7 st_x= StandardScaler()
```


```

8 x_train= st_x.fit_transform(x_train)

1 error = []
2 # Calculating error for K values between 1 and 30
3 for i in range(1, 30):
4     knn = KNeighborsClassifier(n_neighbors=i)
5     knn.fit(x_train, y_train)
6     pred_i = knn.predict(x_test)
7     error.append(np.mean(pred_i != y_test))
8 plt.figure(figsize=(8, 4))
9 plt.plot(range(1, 30), error, color='red', linestyle='dashed', marker='o',
10         markerfacecolor='blue', markersize=10)
11 plt.title('Error Rate K Value')
12 plt.xlabel('K Value')
13 plt.ylabel('Mean Error')
14 print("Minimum error:-",min(error),"at K =",error.index(min(error))+1)
15 plt.show()

```

Minimum error:- 0.36 at K = 7



## ▼ Experiment with various k values and Distance Metrics on UCI Health Dataset

```

1 import pandas as pd
2 import numpy as np
3 from sklearn.neighbors import KNeighborsClassifier
4 from sklearn.metrics import accuracy_score
5 import matplotlib.pyplot as plt
6
7 # Create empty lists to store error, accuracy, and metrics
8 error = []
9 accuracy = []
10 metrics_list = []
11
12 # Create a DataFrame to store results
13 results_df = pd.DataFrame(columns=['K', 'Mean Error', 'Accuracy', 'Metric'])
14
15 # Assuming you have already split your data into x_train, x_test, y_train, and y_test
16
17 # Define the list of distance metrics to iterate over
18 distance_metrics = ['euclidean', 'chebyshev', 'minkowski', 'cosine']
19
20 for metric in distance_metrics:
21     for k in range(1, 16):
22         knn = KNeighborsClassifier(n_neighbors=k, metric=metric)
23         if metric == 'minkowski':
24             knn = KNeighborsClassifier(n_neighbors=k, metric=metric, metric_params={'p': 3})
25         else:
26             knn = KNeighborsClassifier(n_neighbors=k, metric=metric)
27         knn.fit(x_train, y_train)
28         pred_i = knn.predict(x_test)
29
30         # Calculate mean error and accuracy
31         mean_error = np.mean(pred_i != y_test)

```

```

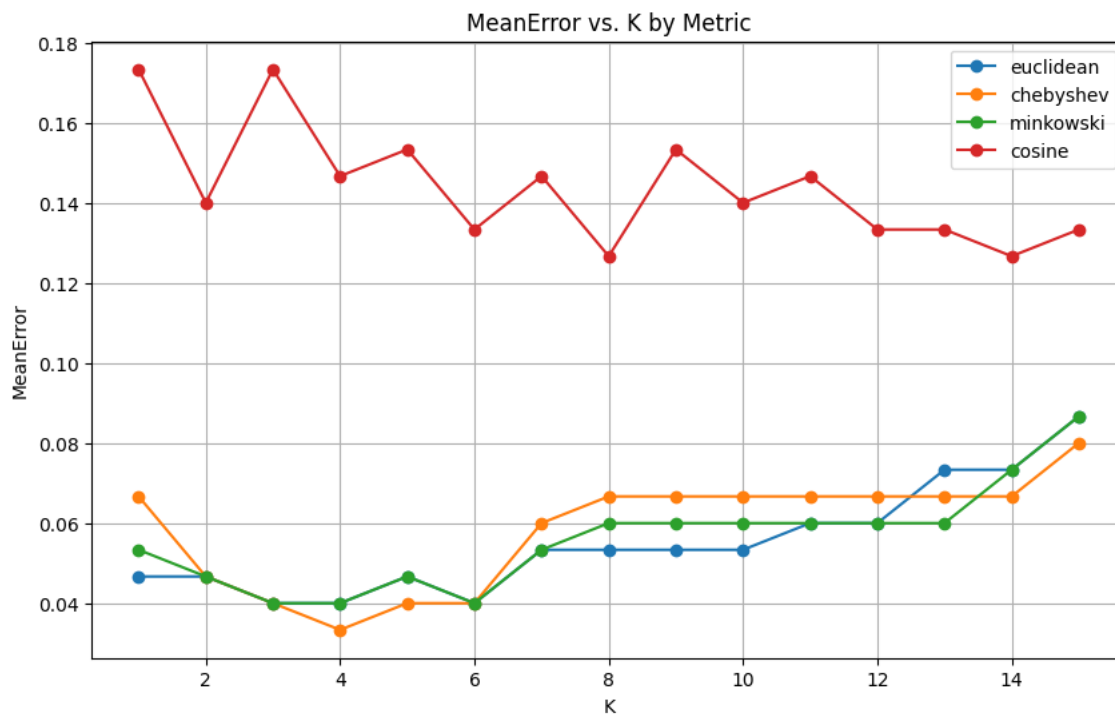
32     acc = accuracy_score(y_test, pred_i)
33
34     # Append values to the lists
35     error.append(mean_error)
36     accuracy.append(acc)
37     metrics_list.append(metric)
38
39     # Append values to the DataFrame
40     results_df = results_df.append({'K': k, 'MeanError': mean_error, 'Accuracy': acc, 'Metric': metric}, ignore_index=True)
41
42 results_df.to_csv("knn.csv", index=False)
43
44
45

```

```

1 import matplotlib.pyplot as plt
2 import pandas as pd
3
4 # Create a figure and axis
5 fig, ax = plt.subplots(figsize=(10, 6))
6 df=results_df
7
8 # Iterate over unique metrics to create lines
9 for metric in df['Metric'].unique():
10     subset = df[df['Metric'] == metric]
11     ax.plot(subset['K'], subset['MeanError'], marker='o', label=metric)
12
13 # Set labels and legend
14 ax.set_xlabel('K')
15 ax.set_ylabel('MeanError')
16 ax.set_title('MeanError vs. K by Metric')
17 ax.legend()
18
19 # Show the plot
20 plt.grid(True)
21 plt.show()
22

```



```

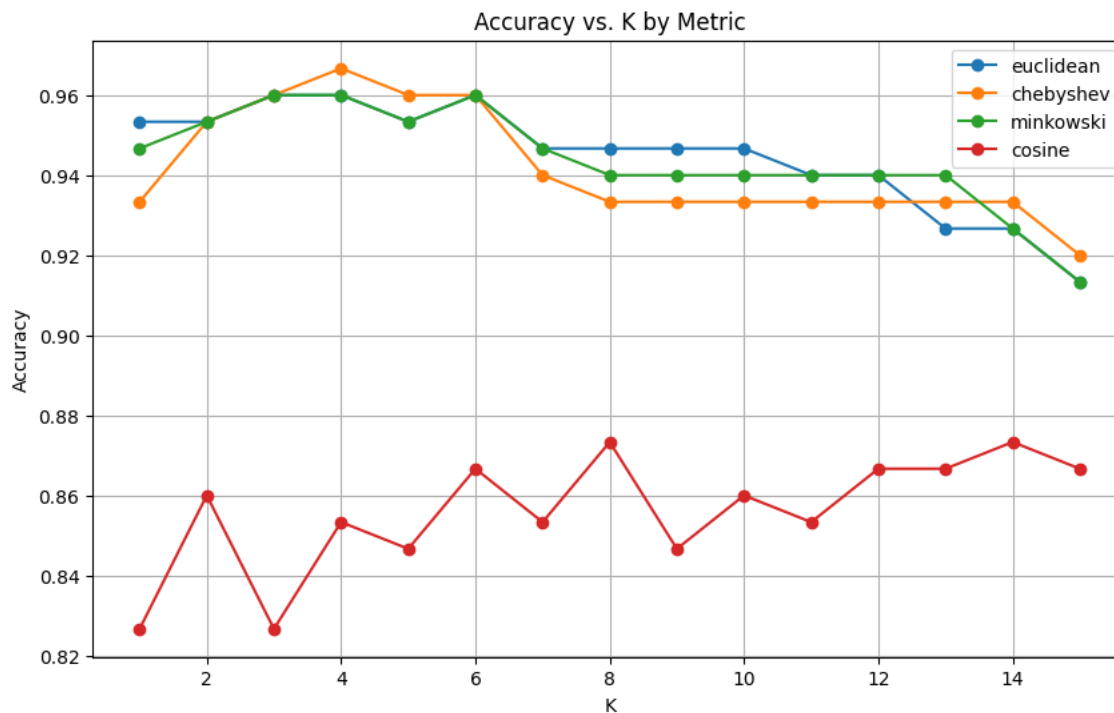
1 import matplotlib.pyplot as plt
2 import pandas as pd
3
4 # Create a figure and axis
5 fig, ax = plt.subplots(figsize=(10, 6))
6 df=results_df
7
8 # Iterate over unique metrics to create lines
9 for metric in df['Metric'].unique():

```

```

10 subset = df[df['Metric'] == metric]
11 ax.plot(subset['K'], subset['Accuracy'], marker='o', label=metric)
12
13 # Set labels and legend
14 ax.set_xlabel('K')
15 ax.set_ylabel('Accuracy')
16 ax.set_title('Accuracy vs. K by Metric')
17 ax.legend()
18
19 # Show the plot
20 plt.grid(True)
21 plt.show()
22

```







## ▼ Create A synthetic dataset to demonstrate the Decision Boundaries at various Iterations

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn.neighbors import KNeighborsClassifier
4 from sklearn.datasets import make_classification
5 from sklearn.model_selection import train_test_split
6 # Generate a synthetic dataset for visualization
7 X, y = make_classification(n_samples=500, n_features=2, n_informative=2, n_redundant=0, random_state=42)
8 x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
9
10 # Calculate error for K values between 1 and 30
11 error = []
12 for k in range(1, 6):
13     knn = KNeighborsClassifier(n_neighbors=k)
14     knn.fit(x_train, y_train)
15     pred_i = knn.predict(x_test)
16     error.append(np.mean(pred_i != y_test))
17
18 # Create a meshgrid for the decision boundary plot
19 h = 0.02 # Step size in the mesh
20 x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
21 y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
22 xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))
23 Z = knn.predict(np.c_[xx.ravel(), yy.ravel()])
24 Z = Z.reshape(xx.shape)
25
26 # Plot the decision boundary for the current K
27 plt.figure(figsize=(10, 5))
28 plt.contourf(xx, yy, Z, cmap=plt.cm.coolwarm, alpha=0.8)

```

```
29 plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.coolwarm)
30 plt.title(f'Decision Boundary for K = {k}')
31 plt.xlabel('Feature 1')
32 plt.ylabel('Feature 2')
33 plt.show()
```




```

1
2 # Plot the error vs. K
3 plt.figure(figsize=(12, 6))
4 plt.plot(range(1, 6), error, color='red', linestyle='dashed', marker='o',

```

```
5     markerfacecolor='blue', markersize=10)
6 plt.title('Error Rate K Value')
7 plt.xlabel('K Value')
8 plt.ylabel('Mean Error')
9 print("Minimum error:", min(error), "at K =", error.index(min(error)) + 1)
```

Minimum error: 0.04 at K = 3



## Summary:

In this project we focused on evaluating the performance of a k-Nearest Neighbors (k-NN) classifier using different distance metrics. The primary objective is to determine the most suitable distance metric and the optimal value of k for a classification task. The code systematically explores a range of k values from 1 to 15 and four different distance metrics: Euclidean, Chebyshev, Minkowski with a specified parameter, and Cosine.

This exercise demonstrates the importance of hyperparameter tuning, specifically the choice of distance metric and the selection of k, in optimizing the performance of k-NN classifiers for classification tasks.

Also the resulting visualizations display the decision boundaries for different K values, helping to understand how the choice of K affects the model's capacity to distinguish between classes. This exercise serves as an informative illustration of the impact of K on k-NN classification performance.